

AEGIS

Demo Pitch — word-for-word script

7 minutes 30 seconds

Flow: citizen → command centre

Counterfeit → Scam → Fraud Graph → Research Lab → Dashboard & GenAI → Services

Highlighted text = the line that must land.

Cue cards at the end of every section.

Before you start

Do this	Why
15 min before: open all four service /health URLs	Free-tier services sleep; waking takes up to 90 seconds. A sleeping service looks broken.
Open every tab once on the dashboard	Warms the caches so nothing loads slowly on stage
Load the scam page and scan one message	Downloads the 3.3 MB on-device model. The aeroplane-mode moment fails without this.
Have a real note ready	For the counterfeit scan
Keep tabs open in this order: counterfeit → scam → dashboard	No fumbling between sections

If the internet fails mid-demo: stay calm and use it. The scam detection runs in the browser — say *"conveniently, this is the part that proves my privacy point: it still works."* Then continue on the dashboard, which holds its last state.

If a service is asleep and a panel is blank: "That's a free-tier server waking up — it takes about a minute." Move on and come back. Do not stand and wait.

On screen: title slide, or the dashboard map

SAY

India registered **1.14 million cybercrime complaints** in 2023 — up sixty percent in a year.

Digital-arrest scams alone took **one thousand seven hundred and seventy-six crore rupees** in nine months.

Now, the police treat these as three separate crimes: the scam call, the mule accounts that move the money, and the fake currency it turns into. Three departments. Three systems. They never talk.

But it isn't three crimes. It's one pipeline, run by one organisation.

Aegis is the first platform that connects all three. Let me show you, starting where a citizen actually stands.

CUE CARD

- 1.14 million complaints, up 60%
- ₹1,776 crore — digital arrest, 9 months
- Police see 3 crimes — 3 departments, never talk
- **"It isn't three crimes. It's one pipeline."**
- Start where the citizen stands

On screen: the counterfeit scanner — **Do:** photograph a real note, live

SAY

A shopkeeper gets a five-hundred rupee note. Is it real? Today he needs a bank machine. Here, he needs a phone.

(scan the note)

Genuine — ninety-nine percent confident. And this is the part that matters: it doesn't just say genuine, **it shows you where it looked.**

That heatmap is the model's attention — the security thread, the portrait, the micro print. If this were fake, those regions light up red, and the officer gets a reason, not just a verdict.

DROP IF SHORT Trained on real photographs of real and counterfeit notes — **ninety-nine point four percent accuracy.** And it remembers every serial it sees: the same serial in two districts means one printing run.

CUE CARD

- Shopkeeper problem → today needs a machine, here a phone
- **Scan the note live**
- Heatmap = *shows where it looked*, not just a verdict
- 0.994 AUC, trained on real photographs
- Serial registry → same serial, two districts = one printing run

Keep this fast. 45 seconds. It's the visual hook, not the main event. If the scan is slow, keep talking over it — do not stand in silence.

On screen: the message checker → then Live Call Shield

SAY

Now the crime that starts it all. This is a real digital-arrest script.

(paste the scam message)

Scam — ninety-nine percent. And notice **why**: authority impersonation, a fake FIR, urgency pressure. Named reasons, not a black-box score. That matters, because this has to stand up in court.

But here is the part I actually want you to remember.

(turn on aeroplane mode)

I've just switched off the network. Watch.

(paste another scam message — it detects)

The AI is running inside the phone. The message never left this browser. We send our servers a verdict — scam, KYC type, district — and nothing else. **We never see what the citizen wrote.**

Same model as the server — identical to eleven decimal places, in under a millisecond.

(switch to Live Call Shield, play the scripted call)

And it works live, during the call. Listen to the risk climb — courier parcel, harmless. "I am Inspector Sharma from the CBI." Now it moves. FIR, digital arrest — **and it interrupts, before the money is asked for.**

Eleven Indian languages. And it catches Hinglish — Hindi typed in English letters — which is how most real scam messages in this country are actually written.

CUE CARD

- Paste digital-arrest script → 99%, **named reasons** (court-usable)
- **AEROPLANE MODE** → paste again → still works
- "The AI runs inside the phone. *We never see what the citizen wrote.*"
- Identical to 11 decimals · 0.8 ms
- Live call: risk climbs → **interrupts before the money ask**
- 11 languages + Hinglish

The aeroplane-mode beat is your best moment. Pause after "watch." Let them see the wifi icon go off. Do not rush it.

On screen: dashboard → Fraud Rings tab — **Do:** inject a ring live

SAY

So the victim pays. Where does the money go?

Not to one account — that's traceable. It's split across dozens belonging to ordinary people paid two thousand rupees for their bank card. Money mules. India has **two and a half million** flagged.

And no single transaction looks wrong. Forty-eight thousand rupees between two normal accounts is normal. **The crime is only visible in the shape of the network.**

(click inject a ring)

I've just created a brand-new laundering ring — accounts that have never existed before. Watch.

Caught. In seconds. It found the whole ring, and it can tell you the role of each account: who collects, who distributes, who is just a link in the chain.

It measures behaviour — money in and out within minutes, round numbers repeatedly, accounts all transacting in a closed circle. Real people don't do that.

Zero point nine nine nine seven accuracy — and **zero false accusations** in evaluation. Every account it named was genuinely in a ring. For something that could freeze an innocent person's account, that number is the one that matters.

And we didn't only test it on our own data. We ran it against **two hundred and sixty-five thousand real Bitcoin wallets** from a published academic dataset. Nought point nine nine four five.

CUE CARD

- Victim pays → money splits across mules → 2.5M flagged in India
- "No single transaction looks wrong — *the crime is in the shape*"
- **Inject a ring live** → caught in seconds
- Behaviour: in-and-out in minutes, round numbers, closed circle
- 0.9997 · **zero false accusations** — the number that matters
- Real benchmark: **265,354 Bitcoin wallets, 0.9945**

On screen: Research Lab tab — **Do:** play the spectral sound

SAY

That's detection working today. Here's what we built to test whether it survives tomorrow.

First problem: criminals deliberately spread accounts across **different banks**. Each bank sees a quarter of the ring and dismisses it as noise. And by law, banks **cannot** share customer data to compare notes.

So we built Ghost Ring. Four banks, each training its own AI, sharing **no customer data** — only an abstract fingerprint, with privacy noise on top.

The best any single bank could see was twenty-three percent of the ring. Fused, **forty-nine percent. Two point two times better** — with zero false links. And the privacy protection cost us **nothing** in accuracy.

DROP IF SHORT Second: we let a simulated criminal **evolve** against our detector for twenty generations. It survives — and where it didn't, we published that too.

Third — and this one you can hear.

(play clean community, then ring)

That's a normal group of accounts. And that's a fraud ring. **A criminal network literally vibrates at a different frequency.**

CUE CARD

- Bridge: "*detection today* → *does it survive tomorrow*"
- Banks can't share data by law → **Ghost Ring**
- Best single bank **23%** → fused **49%** = **2.2×**, zero false links
- Differential privacy at **no accuracy cost**
- Arms race: criminal evolves 20 generations (15 sec, don't linger)
- **Play the sound** — clean vs ring

Discipline here. Ghost Ring gets 45 seconds, the other two get 15 each. The sound does the work — don't explain the mathematics.

On screen: Live Map → Alerts & Analytics (money trail card) → AI case file

SAY

Now everything comes together. This is what the officer sees.

Every scam call, every fake note, every ring — on one map. Where two crime types converge on a district, we mark a coordinated hub: the **infrastructure** is shared, so that isn't coincidence.

And this is the moment I'd like you to take away.

(open the money-trail card)

A victim reported a scam call, and reported paying **forty-nine thousand nine hundred and ninety-nine rupees**. Our fraud graph, working completely separately, knows every transaction between mule accounts.

The system matched them. **That exact amount, arriving minutes after the call, in this account — inside a ring we had already detected.**

Two independent systems, one reading text and one reading bank transactions, pointing at the same crime. That's an account number an officer can freeze today.

DROP IF SHORT One click writes the case file — timeline, hypothesis and next steps, each citing its evidence.

And one thing about the AI, because you'll ask. **The AI never decides anything.** Every verdict here comes from a deterministic model we can audit line by line. The AI only writes the English explanation. Take it out entirely and every verdict is identical — you just lose the readable summary. **That's how we make this court-defensible.**

CUE CARD

- Everything on one map → coordinated hubs = shared infrastructure
- **MONEY TRAIL — the climax**
- Victim paid **₹49,999** → matched into a detected ring's account
- "Two independent systems pointing at the same crime"
- "An account number an officer can freeze *today*"
- One click → AI case file
- **"The AI never decides. It only explains."**

Slow down for the money trail. It's the payoff for everything before it. Pause after "the same crime."

6 · Services & close

6:55 – 7:30

On screen: Modules tab, or an architecture slide

SAY

Last thing. The challenge named three stakeholders, and all three have a working surface.

Citizens get the tools you've seen. Police get this dashboard. Banks get the same models as machine-to-machine APIs, behind a key — so a teller terminal or an AML system calls them directly.

Four services, one validated contract, deployed and running right now.

So — three crimes that are actually one pipeline. Detected at the point of contact, not the point of complaint. **The citizen's message never leaves their phone.** The verdicts are auditable enough for a courtroom. And where our results were weak, we've published that too.

That's Aegis. Thank you.

CUE CARD

- Three stakeholders: citizens · police · **banks via B2B API**
- Four services, one contract, deployed now
- Close: one pipeline · point of contact · *never leaves the phone* · auditable · honest about weaknesses
- **"That's Aegis. Thank you."**

Q&A — the answers to have ready

If they ask	Say
"Does the message go to your server?"	"No. The classifier runs in the browser — same model, verified score-for-score. You saw it work in aeroplane mode."
"What do you send, then?"	"Verdict, scam type, district, timestamp. Not the message. That's all the crime map needs."
"How do you stop AI hallucination?"	"The AI never decides. Every verdict is from a deterministic model. The AI only writes the explanation. Remove it and nothing changes."
"Is this just synthetic data?"	"The counterfeit model is real photographs. The fraud graph was benchmarked on 265,354 real Bitcoin wallets at 0.9945. The scam model includes a public SMS spam corpus."
"What doesn't work?"	"Three things. Investment scams — we catch about half; they read like real financial advice. Live-call speech-to-text goes through the browser, so audio reaches Google. And our synthetic fraud numbers are flattering, which is why we benchmarked on real Bitcoin data."
"Could this be deployed?"	"For SMS yes — Android can read messages, iOS has an on-device SMS filter API for exactly this. For calls, no app on either platform may tap call audio, so that needs telecom-operator integration. Policy constraint, not technical."
"How is this different from Truecaller?"	"Truecaller flags <i>numbers</i> . We classify <i>content</i> — so we work on first contact from an unknown number, and we monitor the call while it's happening."
"Isn't RBI already doing this?"	"MuleHunter flags mule accounts per bank. We built that, plus the layer that joins the ring back to the scam that fed it — and a version that works across banks without them sharing customer data."

Three things never to say:

- "Everything runs on-device" — say "*scam analysis is on-device; speech-to-text is the browser's service*"
- "Our AI decides if it's a scam" — say "*a deterministic model decides; the AI explains*"
- "100% accurate" — quote the real number, and the category where you're weakest

Timing at a glance

Time	Section	The one thing that must land
------	---------	------------------------------

0:00 – 0:35	The problem	"It isn't three crimes. It's one pipeline."
0:35 – 1:20	Counterfeit	The heatmap — it shows you <i>where it looked</i>
1:20 – 3:00	Scam	Aeroplane mode — the message never leaves the phone
3:00 – 4:20	Fraud Graph	Ring injected live, caught in seconds · zero false accusations
4:20 – 5:35	Research Lab	Ghost Ring 2.2x · privacy at no cost · hear the ring
5:35 – 6:55	Dashboard & GenAI	₹49,999 traced into a freezable account
6:55 – 7:30	Services & close	Three stakeholders · "That's Aegis."

If you are running long, cut in this order: (1) the arms race and spectral sound, (2) the counterfeit serial-registry line, (3) the Live Call Shield replay. **Never cut** aeroplane mode or the money trail — they are the demo.